

```
#check that java is installed
!java -version
```

```
openjdk version "11.0.27" 2025-04-15
OpenJDK Runtime Environment (build 11.0.27+6-post-Ubuntu-0ubuntu122.04)
OpenJDK 64-Bit Server VM (build 11.0.27+6-post-Ubuntu-0ubuntu122.04, mixed mode, sharing)
```

```
#install pyspark
!pip install pyspark
```

```
Requirement already satisfied: pyspark in /usr/local/lib/python3.11/dist-packages (3.5.1)
Requirement already satisfied: py4j==0.10.9.7 in /usr/local/lib/python3.11/dist-packages (from pyspark) (0.10.9.7)
```

```
import os
import pandas as pd
from pyspark.sql import SparkSession
from pyspark.sql.types import StructType, StructField, StringType, IntegerType, DoubleType, LongType, TimestampType
from pyspark.sql.functions import spark_partition_id
```

```
spark = SparkSession.builder.appName('AutomotiveData').getOrCreate()

print(f'The Spark version is {spark.version}')
```

```
The Spark version is 3.5.1
```

```
#Table creation for:
#transactions_df (large table)
#countries_df (small table)
```

```
from pyspark.sql import Row
```

```
#Broadcast join
```

```
transactions_data = [
    Row(id=1, countrycode="IN", amount=1000), Row(id=2, countrycode="US", amount=1500), Row(id=3, countrycode="FR", amount=900),
    Row(id=4, countrycode="IN", amount=700), Row(id=5, countrycode="US", amount=1200), Row(id=6, countrycode="FR", amount=1300),
    Row(id=7, countrycode="DE", amount=1100), Row(id=8, countrycode="IN", amount=1400), Row(id=9, countrycode="CA", amount=800),
    Row(id=10, countrycode="US", amount=1700), Row(id=11, countrycode="DE", amount=950), Row(id=12, countrycode="FR", amount=1150),
    Row(id=13, countrycode="IN", amount=1600), Row(id=14, countrycode="CA", amount=1250)
]
transactions_df = spark.createDataFrame(transactions_data)
```

```
countries_data = [Row(countrycode="IN", countryname="India"),
                  Row(countrycode="US", countryname="USA"),
                  Row(countrycode="FR", countryname="France")]
countries_df = spark.createDataFrame(countries_data)
```

```
spark.conf.set("spark.sql.adaptive.enabled", "false")
spark.conf.get("spark.sql.adaptive.enabled")
```

```
'false'
```

```
joined_df = transactions_df.join((countries_df), on="countrycode", how="inner")
joined_df.show()
```

```
+-----+-----+-----+
|countrycode| id|amount|countryname|
+-----+-----+-----+
|          US| 2| 1500|          USA|
|          US| 5| 1200|          USA|
|          US|10| 1700|          USA|
|          IN| 1| 1000|          India|
|          IN| 4|  700|          India|
|          IN| 8| 1400|          India|
|          IN|13| 1600|          India|
|          FR| 3|  900|          France|
|          FR| 6| 1300|          France|
|          FR|12| 1150|          France|
+-----+-----+-----+
```

```
joined_df.explain(True)
```

```

== Parsed Logical Plan ==
'Join UsingJoin(Inner, [countrycode])
:- LogicalRDD [id#0L, countrycode#1, amount#2L], false
+- LogicalRDD [countrycode#6, countryname#7], false

== Analyzed Logical Plan ==
countrycode: string, id: bigint, amount: bigint, countryname: string
Project [countrycode#1, id#0L, amount#2L, countryname#7]
+- Join Inner, (countrycode#1 = countrycode#6)
  :- LogicalRDD [id#0L, countrycode#1, amount#2L], false
  +- LogicalRDD [countrycode#6, countryname#7], false

== Optimized Logical Plan ==
Project [countrycode#1, id#0L, amount#2L, countryname#7]
+- Join Inner, (countrycode#1 = countrycode#6)
  :- Filter isnotnull(countrycode#1)
    : +- LogicalRDD [id#0L, countrycode#1, amount#2L], false
  +- Filter isnotnull(countrycode#6)
    +- LogicalRDD [countrycode#6, countryname#7], false

== Physical Plan ==
*(5) Project [countrycode#1, id#0L, amount#2L, countryname#7]
+- *(5) SortMergeJoin [countrycode#1], [countrycode#6], Inner
  :- *(2) Sort [countrycode#1 ASC NULLS FIRST], false, 0
    : +- Exchange hashpartitioning(countrycode#1, 200), ENSURE_REQUIREMENTS, [plan_id=93]
    :   +- *(1) Filter isnotnull(countrycode#1)
    :     +- *(1) Scan ExistingRDD[id#0L,countrycode#1,amount#2L]
  +- *(4) Sort [countrycode#6 ASC NULLS FIRST], false, 0
    +- Exchange hashpartitioning(countrycode#6, 200), ENSURE_REQUIREMENTS, [plan_id=99]
      +- *(3) Filter isnotnull(countrycode#6)
        +- *(3) Scan ExistingRDD[countrycode#6,countryname#7]

```

```
from pyspark.sql.functions import broadcast
```

```
joined_df_broadcast= transactions_df.join(broadcast(countries_df),on="countrycode",how="inner")
joined_df_broadcast.show()
```

```

+-----+-----+-----+-----+
|countrycode| id|amount|countryname|
+-----+-----+-----+-----+
|          IN| 1| 1000|        India|
|          US| 2| 1500|         USA|
|          FR| 3|  900|       France|
|          IN| 4|  700|        India|
|          US| 5| 1200|         USA|
|          FR| 6| 1300|       France|
|          IN| 8| 1400|        India|
|          US|10| 1700|         USA|
|          FR|12| 1150|       France|
|          IN|13| 1600|        India|
+-----+-----+-----+-----+

```

```
joined_df_broadcast.explain(True)
```

```

== Parsed Logical Plan ==
'Join UsingJoin(Inner, [countrycode])
:- LogicalRDD [id#0L, countrycode#1, amount#2L], false
+- ResolvedHint (strategy=broadcast)
  +- LogicalRDD [countrycode#6, countryname#7], false

== Analyzed Logical Plan ==
countrycode: string, id: bigint, amount: bigint, countryname: string
Project [countrycode#1, id#0L, amount#2L, countryname#7]
+- Join Inner, (countrycode#1 = countrycode#6)
  :- LogicalRDD [id#0L, countrycode#1, amount#2L], false
  +- ResolvedHint (strategy=broadcast)
    +- LogicalRDD [countrycode#6, countryname#7], false

== Optimized Logical Plan ==
Project [countrycode#1, id#0L, amount#2L, countryname#7]
+- Join Inner, (countrycode#1 = countrycode#6), rightHint=(strategy=broadcast)
  :- Filter isnotnull(countrycode#1)
    : +- LogicalRDD [id#0L, countrycode#1, amount#2L], false
  +- Filter isnotnull(countrycode#6)
    +- LogicalRDD [countrycode#6, countryname#7], false

== Physical Plan ==
*(2) Project [countrycode#1, id#0L, amount#2L, countryname#7]
+- *(2) BroadcastHashJoin [countrycode#1], [countrycode#6], Inner, BuildRight, false

```

```

:- *(2) Filter isnonnull(countrycode#1)
: +- *(2) Scan ExistingRDD[id#0L,countrycode#1,amount#2L]
+- BroadcastExchange HashedRelationBroadcastMode(List(input[0, string, false]),false), [plan_id=188]
+- *(1) Filter isnonnull(countrycode#6)
+- *(1) Scan ExistingRDD[countrycode#6,countryname#7]

```

SQL HINTS

```

countries_df.createOrReplaceTempView("countries")
transactions_df.createOrReplaceTempView("transactions")

```

```
spark.sql("select * from countries").show()
```

```

+-----+-----+
|countrycode|countryname|
+-----+-----+
|          IN|        India|
|          US|         USA|
|          FR|        France|
+-----+-----+

```

```

result = spark.sql("""
SELECT
    t.id,
    t.countrycode,
    t.amount,
    c.countryname
FROM
    transactions t
INNER JOIN
    countries c
ON
    t.countrycode = c.countrycode
""")
result.show()

```

```

+---+-----+-----+-----+
| id|countrycode|amount|countryname|
+---+-----+-----+-----+
|  2|          US|  1500|         USA|
|  5|          US|  1200|         USA|
| 10|          US|  1700|         USA|
|  1|          IN|  1000|        India|
|  4|          IN|   700|        India|
|  8|          IN|  1400|        India|
| 13|          IN|  1600|        India|
|  3|          FR|   900|        France|
|  6|          FR|  1300|        France|
| 12|          FR|  1150|        France|
+---+-----+-----+-----+

```

```

result_SQLhint = spark.sql("""
SELECT /*+ BROADCAST(c) */
    t.id,
    t.countrycode,
    t.amount,
    c.countryname
FROM
    transactions t
INNER JOIN
    countries c
ON
    t.countrycode = c.countrycode
""")
result_SQLhint.show()

```

```

+---+-----+-----+-----+
| id|countrycode|amount|countryname|
+---+-----+-----+-----+
|  1|          IN|  1000|        India|
|  2|          US|  1500|         USA|
|  3|          FR|   900|        France|

```

	4		IN		700		India	
	5		US		1200		USA	
	6		FR		1300		France	
	8		IN		1400		India	
	10		US		1700		USA	
	12		FR		1150		France	
	13		IN		1600		India	
+---+-----+-----+-----+								