

```
#check that java is installed
!java -version
```

```
OpenJDK version "11.0.27" 2025-04-15
OpenJDK Runtime Environment (build 11.0.27+6-post-Ubuntu-0ubuntu122.04)
OpenJDK 64-Bit Server VM (build 11.0.27+6-post-Ubuntu-0ubuntu122.04, mixed mode, sharing)
```

```
#install pyspark
!pip install pyspark
```

```
Requirement already satisfied: pyspark in /usr/local/lib/python3.11/dist-packages (3.5.1)
Requirement already satisfied: py4j==0.10.9.7 in /usr/local/lib/python3.11/dist-packages (from pyspark) (0.10.9.7)
```

```
import os
import pandas as pd
from pyspark.sql import SparkSession
from pyspark.sql.types import StructType, StructField, StringType, IntegerType, DoubleType, LongType, TimestampType
from pyspark.sql.functions import spark_partition_id
```

```
!curl -O https://raw.githubusercontent.com/Apoorva-888/Spark-Optimization/main/BigMart_Sales.csv
```

```
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
   100    849k   100    849k    0     0  2019k      0  --:--:--  --:--:--  --:--:--  2017k
```

[+ Code](#)
[+ Text](#)

```
spark = SparkSession.builder.appName('AutomotiveData').getOrCreate()
```

```
print(f'The Spark version is {spark.version}')
```

```
The Spark version is 3.5.1
```

```
spark.conf.set("spark.sql.adaptive.enabled", "false")
spark.conf.get("spark.sql.adaptive.enabled")
```

```
'false'
```

```
spark.conf.get("spark.sql.adaptive.enabled")
```

```
'false'
```

```
df = spark.read.format("csv") \
    .option("inferSchema", True) \
    .option("header", True) \
    .load("BigMart_Sales.csv")
```

```
df.show(5)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
|Item_Identifier|Item_Weight|Item_Fat_Content|Item_Visibility|Item_Type|Item_MRP|Outlet_Identifier|Outlet_Establishment_Year|
+-----+-----+-----+-----+-----+-----+-----+-----+
|FDA15|9.3|Low Fat|0.016047301|Dairy|249.8092|OUT049|1999|
|DRC01|5.92|Regular|0.019278216|Soft Drinks|48.2692|OUT018|2009|
|FDN15|17.5|Low Fat|0.016760075|Meat|141.618|OUT049|1999|
|FDX07|19.2|Regular|0.0|Fruits and Vegeta...|182.095|OUT010|1998|
|NCD19|8.93|Low Fat|0.0|Household|53.8614|OUT013|1987|
+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows
```

```
df.printSchema()
```

```
root
 |-- Item_Identifier: string (nullable = true)
 |-- Item_Weight: double (nullable = true)
 |-- Item_Fat_Content: string (nullable = true)
 |-- Item_Visibility: double (nullable = true)
 |-- Item_Type: string (nullable = true)
 |-- Item_MRP: double (nullable = true)
 |-- Outlet_Identifier: string (nullable = true)
 |-- Outlet_Establishment_Year: integer (nullable = true)
 |-- Outlet_Size: string (nullable = true)
```

```
-- Outlet_Location_Type: string (nullable = true)
-- Outlet_Type: string (nullable = true)
-- Item_Outlet_Sales: double (nullable = true)
```

```
!ls -lh BigMart_Sales.csv
```

```
-rw-r--r-- 1 root root 850K Jul 17 14:21 BigMart_Sales.csv
```

```
spark.conf.get("spark.sql.files.maxPartitionBytes")
```

```
'134217728b'
```

```
df.rdd.getNumPartitions()
```

```
1
```

Changing default partition to 128kb

Start coding or [generate](#) with AI.

```
spark.conf.set("spark.sql.files.maxPartitionBytes", 131072)
df = spark.read.format("csv").option("inferSchema", True).option("header", True).load("BigMart_Sales.csv")
df.rdd.getNumPartitions()
```

```
7
```

Changing default partition to 128MB

```
spark.conf.set("spark.sql.files.maxPartitionBytes", 1024 * 1024 * 1)
df = spark.read.format("csv").option("inferSchema", True).option("header", True).load("BigMart_Sales.csv")
df.rdd.getNumPartitions()
```

```
1
```

Repartition

```
df=df.repartition(10)
df.rdd.getNumPartitions()
```

```
10
```

```
#Write a DataFrame as Parquet in append mode in Google Colab.
df.write.format("parquet").mode("append").option("path", "/content/parquetWrite").save()
```

```
#After Saving: You can verify the files were written:
!ls -lh /content/parquetWrite
```

```
total 324K
-rw-r--r-- 1 root root 32K Jul 17 14:22 part-00000-383227c2-52d8-4466-afe5-868d48f6a786-c000.snappy.parquet
-rw-r--r-- 1 root root 31K Jul 17 14:22 part-00001-383227c2-52d8-4466-afe5-868d48f6a786-c000.snappy.parquet
-rw-r--r-- 1 root root 31K Jul 17 14:22 part-00002-383227c2-52d8-4466-afe5-868d48f6a786-c000.snappy.parquet
-rw-r--r-- 1 root root 32K Jul 17 14:22 part-00003-383227c2-52d8-4466-afe5-868d48f6a786-c000.snappy.parquet
-rw-r--r-- 1 root root 31K Jul 17 14:22 part-00004-383227c2-52d8-4466-afe5-868d48f6a786-c000.snappy.parquet
-rw-r--r-- 1 root root 33K Jul 17 14:22 part-00005-383227c2-52d8-4466-afe5-868d48f6a786-c000.snappy.parquet
-rw-r--r-- 1 root root 32K Jul 17 14:22 part-00006-383227c2-52d8-4466-afe5-868d48f6a786-c000.snappy.parquet
-rw-r--r-- 1 root root 32K Jul 17 14:22 part-00007-383227c2-52d8-4466-afe5-868d48f6a786-c000.snappy.parquet
-rw-r--r-- 1 root root 32K Jul 17 14:22 part-00008-383227c2-52d8-4466-afe5-868d48f6a786-c000.snappy.parquet
-rw-r--r-- 1 root root 31K Jul 17 14:22 part-00009-383227c2-52d8-4466-afe5-868d48f6a786-c000.snappy.parquet
-rw-r--r-- 1 root root  0 Jul 17 14:22 _SUCCESS
```

```
#Reading It Back:
df2 = spark.read.parquet("/content/parquetWrite")
df2.show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
|Item_Identifier|Item_Weight|Item_Fat_Content|Item_Visibility|Item_Type|Item_MRP|Outlet_Identifier|Outlet_Establishment_Year|
+-----+-----+-----+-----+-----+-----+-----+-----+
|      FDU21      |      NULL      |      Regular      |      0.134327613      |Snack Foods| 35.0558|      OUT019      |      1985      |
```

	FDN56	5.46	Regular	0.106968096	Fruits and Vegeta...	142.6786	OUT013	1987
	NCJ17	7.68	Low Fat	0.153177941	Health and Hygiene	85.2224	OUT018	2009
	FDG08	NULL	Regular	0.289522833	Fruits and Vegeta...	172.0764	OUT019	1985
	NCZ41	19.85	Low Fat	0.064409056	Health and Hygiene	126.7704	OUT035	2004
	FDI60	7.22	Regular	0.0383808	Baking Goods	62.351	OUT049	1999
	FDR36	6.715	reg	0.122274118	Baking Goods	40.3454	OUT017	2007
	FDZ12	9.17	Low Fat	0.102978817	Baking Goods	144.947	OUT046	1997
	NCM07	9.1	Low Fat	0.055382616	Household	115.0518	OUT013	1987
	DRG39	14.15	Low Fat	0.042352822	Dairy	51.6982	OUT018	2009
	FDR19	13.5	Regular	0.160624116	Fruits and Vegeta...	147.0102	OUT017	2007
	FDK56	9.695	Low Fat	0.130261652	Fruits and Vegeta...	185.2898	OUT045	2002
	NCM07	NULL	Low Fat	0.03976832	Others	83.9908	OUT027	1985
	NCI30	20.25	Low Fat	0.059055045	Household	247.346	OUT045	2002
	FDZ35	9.6	Regular	0.02236923	Breads	104.799	OUT018	2009
	FDB14	20.25	Regular	0.103142373	Canned	94.612	OUT018	2009
	NCN05	8.235	Low Fat	0.014541462	Health and Hygiene	184.495	OUT017	2007
	FDZ38	17.6	LF	0.008001018	Dairy	170.4422	OUT046	1997
	FDM38	5.885	Regular	0.092694107	Canned	53.6982	OUT013	1987
	DRK13	11.8	Low Fat	0.115346634	Soft Drinks	200.2084	OUT049	1999

only showing top 20 rows

```
from pyspark.sql.functions import spark_partition_id
df_with_partition = df.withColumn("partition_id", spark_partition_id())
df_with_partition.show(truncate=False) # Shows rows with partition info
df_with_partition.groupBy("partition_id").count().show()
```

Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Type
DRG48	5.78	Low Fat	0.014555066	Soft Drinks	145.2102	OUT046	1997	Sr
FDE24	14.85	Low Fat	0.09344495	Baking Goods	141.0812	OUT035	2004	Sr
NCP18	12.15	Low Fat	0.028714747	Household	151.9708	OUT018	2009	Me
FDX27	20.7	Regular	0.114581955	Dairy	94.3436	OUT018	2009	Me
FDU02	13.35	Low Fat	0.102670882	Dairy	228.6352	OUT049	1999	Me
FDW57	8.31	Regular	0.115911972	Snack Foods	177.3028	OUT045	2002	NL
FDZ01	8.975	Regular	0.009057132	Canned	104.099	OUT035	2004	Sr
FDL39	16.1	Regular	0.063689583	Dairy	181.9318	OUT017	2007	NL
FDL43	10.1	Low Fat	0.027064381	Meat	76.367	OUT046	1997	Sr
FDP23	NULL	Low Fat	0.035414528	Breads	218.2166	OUT027	1985	Me
FDL15	17.85	Low Fat	0.046824729	Meat	153.6682	OUT018	2009	Me
FDB23	19.2	Regular	0.0	Starchy Foods	223.8062	OUT045	2002	NL
FDY39	NULL	Regular	0.08234117	Meat	185.7608	OUT019	1985	Sr
FDG05	11.0	Regular	0.0	Frozen Foods	155.263	OUT049	1999	Me
NCS29	9.0	Low Fat	0.069653585	Health and Hygiene	266.2884	OUT049	1999	Me
FDT01	13.65	Regular	0.184454044	Canned	211.4902	OUT049	1999	Me
NCJ05	18.7	Low Fat	0.046348967	Health and Hygiene	153.6682	OUT017	2007	NL
NCX42	6.36	Low Fat	0.00599072	Household	163.6526	OUT045	2002	NL
FDF10	15.5	reg	0.156797787	Snack Foods	148.6418	OUT013	1987	Hi
FDZ15	13.1	Low Fat	0.0	Dairy	117.8782	OUT035	2004	Sr

only showing top 20 rows

partition_id	count
1	853
6	852
3	852
5	852
9	853
4	852
8	852
7	852
2	852
0	853

## ✓ coalesce and repartition

```
df2 = spark.read.parquet("/content/parquetWrite")
df2.show()
```

Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Type
-----------------	-------------	------------------	-----------------	-----------	----------	-------------------	---------------------------	-------------

FDU21	NULL	Regular	0.134327613	Snack Foods	35.0558	OUT019	1985
FDN56	5.46	Regular	0.106968096	Fruits and Vegeta...	142.6786	OUT013	1987
NCJ17	7.68	Low Fat	0.153177941	Health and Hygiene	85.2224	OUT018	2009
FDG08	NULL	Regular	0.289522833	Fruits and Vegeta...	172.0764	OUT019	1985
NCZ41	19.85	Low Fat	0.064409056	Health and Hygiene	126.7704	OUT035	2004
FDI60	7.22	Regular	0.0383808	Baking Goods	62.351	OUT049	1999
FDR36	6.715	reg	0.122274118	Baking Goods	40.3454	OUT017	2007
FDZ12	9.17	Low Fat	0.102978817	Baking Goods	144.947	OUT046	1997
NCD07	9.1	Low Fat	0.055382616	Household	115.0518	OUT013	1987
DRG39	14.15	Low Fat	0.042352822	Dairy	51.6982	OUT018	2009
FDR19	13.5	Regular	0.160624116	Fruits and Vegeta...	147.0102	OUT017	2007
FDK56	9.695	Low Fat	0.130261652	Fruits and Vegeta...	185.2898	OUT045	2002
NCM07	NULL	Low Fat	0.03976832	Others	83.9908	OUT027	1985
NCI30	20.25	Low Fat	0.059055045	Household	247.346	OUT045	2002
FDZ35	9.6	Regular	0.02236923	Breads	104.799	OUT018	2009
FDB14	20.25	Regular	0.103142373	Canned	94.612	OUT018	2009
NCN05	8.235	Low Fat	0.014541462	Health and Hygiene	184.495	OUT017	2007
FDZ38	17.6	LF	0.008001018	Dairy	170.4422	OUT046	1997
FDM38	5.885	Regular	0.092694107	Canned	53.6982	OUT013	1987
DRK13	11.8	Low Fat	0.115346634	Soft Drinks	200.2084	OUT049	1999

only showing top 20 rows

```
from pyspark.sql.functions import *
df3=df2.filter(col("Outlet_Location_Type")== 'Tier 1')
```

```
from pyspark.sql.functions import *
df4=df2.filter(col("Outlet_Location_Type")== 'Tier 2')
```

```
df2_cache = spark.read.parquet("/content/parquetWrite").cache()
from pyspark.sql.functions import *
df42=df2_cache.filter(col("Outlet_Location_Type")== 'Tier 2')
df42.show()
```

Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year
NCZ41	19.85	Low Fat	0.064409056	Health and Hygiene	126.7704	OUT035	2004
FDR36	6.715	reg	0.122274118	Baking Goods	40.3454	OUT017	2007
FDR19	13.5	Regular	0.160624116	Fruits and Vegeta...	147.0102	OUT017	2007
FDK56	9.695	Low Fat	0.130261652	Fruits and Vegeta...	185.2898	OUT045	2002
NCI30	20.25	Low Fat	0.059055045	Household	247.346	OUT045	2002
NCN05	8.235	Low Fat	0.014541462	Health and Hygiene	184.495	OUT017	2007
FDU02	13.35	Low Fat	0.1027194	Dairy	228.8352	OUT045	2002
DRI13	15.35	Low Fat	0.020441939	Soft Drinks	216.4508	OUT017	2007
FDO08	11.1	Regular	0.054079556	Fruits and Vegeta...	165.9526	OUT017	2007
FDJ09	15.0	low fat	0.0	Snack Foods	47.2744	OUT035	2004
FDB26	14.0	reg	0.031444357	Canned	53.764	OUT017	2007
FDB21	7.475	Low Fat	0.149360698	Fruits and Vegeta...	243.4854	OUT017	2007
NCL41	12.35	Low Fat	0.041973712	Health and Hygiene	35.7216	OUT017	2007
NCL31	7.39	Low Fat	0.120524922	Others	142.247	OUT045	2002
FDT40	5.985	Low Fat	0.0	Frozen Foods	125.2678	OUT035	2004
NCN55	14.6	Low Fat	0.059827007	Others	239.2538	OUT017	2007
FDB51	6.92	LF	0.038671588	Dairy	64.2852	OUT017	2007
DRG15	6.13	Low Fat	0.076891526	Dairy	61.5536	OUT045	2002
FDF22	6.865	Low Fat	0.056819936	Snack Foods	212.6218	OUT035	2004
FDX50	20.1	Low Fat	0.075049323	Dairy	110.4228	OUT017	2007

only showing top 20 rows

```
df2_cache.unpersist()
```

```
DataFrame[Item_Identifier: string, Item_Weight: double, Item_Fat_Content: string, Item_Visibility: double, Item_Type: string, Item_MRP: double, Outlet_Identifier: string, Outlet_Establishment_Year: int, Outlet_Size: string, Outlet_Location_Type: string, Outlet_Type: string, Item_Outlet_Sales: double]
```

```
from pyspark import StorageLevel
```

```
df3.persist(StorageLevel.MEMORY_AND_DISK)
```

```
DataFrame[Item_Identifier: string, Item_Weight: double, Item_Fat_Content: string, Item_Visibility: double, Item_Type: string, Item_MRP: double, Outlet_Identifier: string, Outlet_Establishment_Year: int, Outlet_Size: string, Outlet_Location_Type: string, Outlet_Type: string, Item_Outlet_Sales: double]
```

```
# Default: MEMORY_AND_DISK (if memory full, spills to disk)
df3.persist(StorageLevel.MEMORY_AND_DISK)

# Other examples:
# df.persist(StorageLevel.MEMORY_ONLY)      # Stores only in memory, errors if not enough space
# df.persist(StorageLevel.DISK_ONLY)        # Skips memory, directly caches to disk
# df.persist(StorageLevel.MEMORY_AND_DISK_SER) # Serialized format, less memory, more CPU

df3.count() # triggers cache
```

 2388

Start coding or [generate](#) with AI.