

```
#check that java is installed
!java -version
```

```
openjdk version "11.0.27" 2025-04-15
OpenJDK Runtime Environment (build 11.0.27+6-post-Ubuntu-0ubuntu122.04)
OpenJDK 64-Bit Server VM (build 11.0.27+6-post-Ubuntu-0ubuntu122.04, mixed mode, sharing)
```

```
#install pyspark
!pip install pyspark
```

```
Requirement already satisfied: pyspark in /usr/local/lib/python3.11/dist-packages (3.5.1)
Requirement already satisfied: py4j==0.10.9.7 in /usr/local/lib/python3.11/dist-packages (from pyspark) (0.10.9.7)
```

```
import os
import pandas as pd
from pyspark.sql import SparkSession
from pyspark.sql.types import StructType, StructField, StringType, IntegerType, DoubleType, LongType, TimestampType
from pyspark.sql.functions import spark_partition_id
```

```
!curl -O https://raw.githubusercontent.com/Apoorva-888/Spark-Optimization/main/BigMart\_Sales.csv
```

```
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total      Spent      Left   Speed
100  849k  100  849k    0     0  4829k      0 --:--:-- --:--:-- --:--:-- 4824k
```

```
spark = SparkSession.builder.appName('AutomotiveData').getOrCreate()

print(f'The Spark version is {spark.version}')
```

```
The Spark version is 3.5.1
```

```
spark.conf.set("spark.sql.adaptive.enabled", "false")
spark.conf.get("spark.sql.adaptive.enabled")
```

```
'false'
```

```
spark.conf.set("spark.sql.optimizer.dynamicPartitionPruning.enabled", "false")
spark.conf.get("spark.sql.optimizer.dynamicPartitionPruning.enabled")
```

```
'false'
```

```
spark.conf.set("spark.sql.autoBroadcastJoinThreshold", -1)
spark.conf.get("spark.sql.autoBroadcastJoinThreshold")
```

```
'-1'
```

```
df = spark.read.format("csv") \
    .option("inferSchema", True) \
    .option("header", True) \
    .load("BigMart_Sales.csv")
```

```
df.show(5)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
|Item_Identifier|Item_Weight|Item_Fat_Content|Item_Visibility|Item_Type|Item_MRP|Outlet_Identifier|Outlet_Establishment_Year|
+-----+-----+-----+-----+-----+-----+-----+-----+
|      FDA15|      9.3|      Low Fat|    0.016047301|      Dairy|249.8092|      OUT049|      1999|
|      DRC01|      5.92|      Regular|    0.019278216|    Soft Drinks|48.2692|      OUT018|      2009|
|      FDN15|      17.5|      Low Fat|    0.016760075|      Meat|141.618|      OUT049|      1999|
|      FDX07|      19.2|      Regular|      0.0|Fruits and Vegeta...|182.095|      OUT010|      1998|
|      NCD19|      8.93|      Low Fat|      0.0|      Household|53.8614|      OUT013|      1987|
+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows
```

```
df.printSchema()
```

```
root
 |-- Item_Identifier: string (nullable = true)
```

```
-- Item_Weight: double (nullable = true)
-- Item_Fat_Content: string (nullable = true)
-- Item_Visibility: double (nullable = true)
-- Item_Type: string (nullable = true)
-- Item_MRP: double (nullable = true)
-- Outlet_Identifier: string (nullable = true)
-- Outlet_Establishment_Year: integer (nullable = true)
-- Outlet_Size: string (nullable = true)
-- Outlet_Location_Type: string (nullable = true)
-- Outlet_Type: string (nullable = true)
-- Item_Outlet_Sales: double (nullable = true)
```

```
lls -lh BigMart_Sales.csv
```

```
-rw-r--r-- 1 root root 850K Jul 21 13:29 BigMart_Sales.csv
```

```
spark.conf.get("spark.sql.files.maxPartitionBytes")
```

```
'134217728b'
```

```
df.rdd.getNumPartitions()
```

```
1
```

✓ Preparing partitioned data

```
df_dpp_partitioned = df.write.format("parquet") \
    .partitionBy("Outlet_Type") \
    .option("path", "dpp_partitioned_BigMart_Sales") \
    .save()
```

```
df_dpp_partitioned = spark.read.parquet("dpp_partitioned_BigMart_Sales")
df_dpp_partitioned.show(5)
```

```
only showing top 5 rows
```

Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Sales
FDA15	9.3	Low Fat	0.016047301	Dairy	249.8092	OUT049	1999	Med
FDN15	17.5	Low Fat	0.016760075	Meat	141.618	OUT049	1999	Med
NCD19	8.93	Low Fat	0.0	Household	53.8614	OUT013	1987	F
FDO10	13.65	Regular	0.012741089	Snack Foods	57.6588	OUT013	1987	F
FDH17	16.2	Regular	0.016687114	Frozen Foods	96.9726	OUT045	2002	N

✓ Preparing Non partitioned data

```
df_dpp_non_partitioned= df.write.format("parquet") \
    .mode("append") \
    .option("path", "/content/dpp_non_partitioned_BigMart_Sales") \
    .save()
```


```
df_dpp_non_partitioned = spark.read.parquet("dpp_non_partitioned_BigMart_Sales")
df_dpp_non_partitioned.show(5)
```

```
only showing top 5 rows
```

Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Sales
FDA15	9.3	Low Fat	0.016047301	Dairy	249.8092	OUT049	1999	Med
DRC01	5.92	Regular	0.019278216	Soft Drinks	48.2692	OUT018	2009	Med
FDN15	17.5	Low Fat	0.016760075	Meat	141.618	OUT049	1999	Med
FDX07	19.2	Regular	0.0	Fruits and Vegetables	182.095	OUT010	1998	F
NCD19	8.93	Low Fat	0.0	Household	53.8614	OUT013	1987	F


only showing top 5 rows

```
import os
os.listdir("/content")
!find /content -name "*BigMart*"
```

 /content/BigMart_Sales.csv
 /content/dpp_non_partitioned_BigMart_Sales
 /content/dpp_partitioned_BigMart_Sales

```
#!rm -r dpp_partitioned_BigMart_Sales
```

```
!ls -R /content
```

 /content:
 BigMart_Sales.csv dpp_partitioned_BigMart_Sales
 dpp_non_partitioned_BigMart_Sales sample_data

/content/dpp_non_partitioned_BigMart_Sales:
 part-00000-aa477b3a-cccc-4bc7-87ca-9f4d0e40c424-c000.snappy.parquet _SUCCESS

/content/dpp_partitioned_BigMart_Sales:
 'Outlet_Type=Grocery Store' 'Outlet_Type=Supermarket Type2' _SUCCESS
 'Outlet_Type=Supermarket Type1' 'Outlet_Type=Supermarket Type3'

'/content/dpp_partitioned_BigMart_Sales/Outlet_Type=Grocery Store':
 part-00000-f41d5f25-07eb-4d91-9f01-8115194828b4.c000.snappy.parquet

'/content/dpp_partitioned_BigMart_Sales/Outlet_Type=Supermarket Type1':
 part-00000-f41d5f25-07eb-4d91-9f01-8115194828b4.c000.snappy.parquet

'/content/dpp_partitioned_BigMart_Sales/Outlet_Type=Supermarket Type2':
 part-00000-f41d5f25-07eb-4d91-9f01-8115194828b4.c000.snappy.parquet

'/content/dpp_partitioned_BigMart_Sales/Outlet_Type=Supermarket Type3':
 part-00000-f41d5f25-07eb-4d91-9f01-8115194828b4.c000.snappy.parquet

/content/sample_data:
 anscombe.json mnist_test.csv
 california_housing_test.csv mnist_train_small.csv
 california_housing_train.csv README.md

Double-click (or enter) to edit

```
##Non-partitioned folder
/content/dpp_non_partitioned_BigMart_Sales/
Contains: 1 .parquet file
```

```
##Partitioned folder
/content/dpp_partitioned_BigMart_Sales/
Contains: 4 folders (based on Outlet_Type), each with 1 .parquet file
```

```
# For partitioned folder (nested subfolders):
```

```
!find /content/dpp_partitioned_BigMart_Sales -name "*.parquet" | wc -l
```

 4

```
#For non-partitioned folder (flat):
```

```
!ls /content/dpp_non_partitioned_BigMart_Sales/*.parquet | wc -l
```

 1


Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

```
from pyspark.sql.functions import *
df_joined = df_dpp_partitioned.join(df_dpp_non_partitioned.filter(col("Outlet_Type")=="Grocery Store"), on="Item_Identifier", how="inner")

df_joined.show()
```




Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_S
DRA24	19.35	Regular	0.040154087	Soft Drinks	164.6868	OUT017	2007	N
DRA24	19.35	Regular	0.040154087	Soft Drinks	164.6868	OUT017	2007	N
DRA24	19.35	Regular	0.039920687	Soft Drinks	163.3868	OUT035	2004	Sr
DRA24	19.35	Regular	0.039920687	Soft Drinks	163.3868	OUT035	2004	Sr
DRA24	19.35	Regular	0.039990314	Soft Drinks	165.0868	OUT049	1999	Med
DRA24	19.35	Regular	0.039990314	Soft Drinks	165.0868	OUT049	1999	Med
DRA24	19.35	Regular	0.039895009	Soft Drinks	162.4868	OUT013	1987	F
DRA24	19.35	Regular	0.039895009	Soft Drinks	162.4868	OUT013	1987	F
DRA24	NULL	Regular	0.069909188	Soft Drinks	163.2868	OUT019	1985	Sr
DRA24	NULL	Regular	0.069909188	Soft Drinks	163.2868	OUT019	1985	Sr
DRA24	19.35	Regular	0.066831682	Soft Drinks	163.8868	OUT010	1998	N
DRA24	19.35	Regular	0.066831682	Soft Drinks	163.8868	OUT010	1998	N
DRA24	NULL	Regular	0.039734882	Soft Drinks	165.7868	OUT027	1985	Med
DRA24	NULL	Regular	0.039734882	Soft Drinks	165.7868	OUT027	1985	Med
FDO11	8.0	Regular	0.030311951	Breads	247.4092	OUT049	1999	Med
FDO11	8.0	Regular	0.030264897	Breads	250.3092	OUT046	1997	Sr
FDO11	8.0	Regular	0.050657232	Breads	249.9092	OUT010	1998	N
FDO11	NULL	Regular	0.030118338	Breads	248.8092	OUT027	1985	Med
FDU24	6.78	Regular	0.140955857	Baking Goods	92.212	OUT017	2007	N
FDU24	6.78	Regular	0.0	Baking Goods	94.012	OUT013	1987	F


only showing top 20 rows

```
#Show joined result and files read

print("Files read during join:", len(df_joined.inputFiles()))
```

 Files read during join: 5

```
df_joined.explain(mode="formatted")
```



```
(1) Scan parquet
Output [12]: [Item_Identifier#1041, Item_Weight#1042, Item_Fat_Content#1043, Item_Visibility#1044, Item_Type#1045, Item_MRP#1046, Outl
Batched: true
Location: InMemoryFileIndex [file:/content/dpp_partitioned_BigMart_Sales]
PushedFilters: [IsNull(Item_Identifier)]
ReadSchema: struct<Item_Identifier:string,Item_Weight:double,Item_Fat_Content:string,Item_Visibility:double,Item_Type:string,Item_MRP:

(2) ColumnarToRow [codegen id : 1]
Input [12]: [Item_Identifier#1041, Item_Weight#1042, Item_Fat_Content#1043, Item_Visibility#1044, Item_Type#1045, Item_MRP#1046, Outle

(3) Filter [codegen id : 1]
Input [12]: [Item_Identifier#1041, Item_Weight#1042, Item_Fat_Content#1043, Item_Visibility#1044, Item_Type#1045, Item_MRP#1046, Outle
Condition : isnotnull(Item_Identifier#1041)

(4) Exchange
Input [12]: [Item_Identifier#1041, Item_Weight#1042, Item_Fat_Content#1043, Item_Visibility#1044, Item_Type#1045, Item_MRP#1046, Outle
Arguments: hashpartitioning(Item_Identifier#1041, 200), ENSURE_REQUIREMENTS, [plan_id=564]

(5) Sort [codegen id : 2]
Input [12]: [Item_Identifier#1041, Item_Weight#1042, Item_Fat_Content#1043, Item_Visibility#1044, Item_Type#1045, Item_MRP#1046, Outle
Arguments: [Item_Identifier#1041 ASC NULLS FIRST], false, 0
```

```
(9) Exchange
Input [12]: [Item_Identifier#942, Item_Weight#943, Item_Fat_Content#944, Item_Visibility#945, Item_Type#946, Item_MRP#947, Outlet_Identifier#948]
Arguments: hashpartitioning(Item_Identifier#942, 200), ENSURE_REQUIREMENTS, [plan_id=573]

(10) Sort [codegen id : 4]
Input [12]: [Item_Identifier#942, Item_Weight#943, Item_Fat_Content#944, Item_Visibility#945, Item_Type#946, Item_MRP#947, Outlet_Identifier#948]
Arguments: [Item_Identifier#942 ASC NULLS FIRST], false, 0

(11) SortMergeJoin [codegen id : 5]
Left keys [1]: [Item_Identifier#1041]
Right keys [1]: [Item_Identifier#942]
Join type: Inner
Join condition: None

(12) Project [codegen id : 5]
Output [23]: [Item_Identifier#1041, Item_Weight#1042, Item_Fat_Content#1043, Item_Visibility#1044, Item_Type#1045, Item_MRP#1046, Outlet_Identifier#1047]
Input [24]: [Item_Identifier#1041, Item_Weight#1042, Item_Fat_Content#1043, Item_Visibility#1044, Item_Type#1045, Item_MRP#1046, Outlet_Identifier#1047]
```

```
#Show joined result and files read

print("Files read during join:", len(df2_joined.inputFiles()))
```

Files read during join: 5

Start coding or generate with AI.

Start coding or generate with AI.

Start coding or generate with AI.

Start coding or generate with AI.

Start coding or generate with AI.

DDP concept

```
# Turn off Adaptive Query Execution
spark.conf.set("spark.sql.adaptive.enabled", "false")
print("Adaptive Query Execution enabled:", spark.conf.get("spark.sql.adaptive.enabled"))

# Enable Dynamic Partition Pruning
spark.conf.set("spark.sql.optimizer.dynamicPartitionPruning.enabled", "true")
print("Dynamic Partition Pruning enabled:", spark.conf.get("spark.sql.optimizer.dynamicPartitionPruning.enabled"))

# Disable Auto Broadcast Join
spark.conf.set("spark.sql.autoBroadcastJoinThreshold", -1)
print("Auto Broadcast Join Threshold:", spark.conf.get("spark.sql.autoBroadcastJoinThreshold"))
```

Adaptive Query Execution enabled: false
Dynamic Partition Pruning enabled: true
Auto Broadcast Join Threshold: -1

```
df2 = spark.read.format("csv").option("inferSchema", True).option("header", True).load("BigMart_Sales.csv")

df2_dpp_partitioned = df2.write.format("parquet").partitionBy("Item_Identifier").option("path", "dpp2_partitioned_BigMart_Sales").save()
df2_dpp_partitioned = spark.read.parquet("dpp2_partitioned_BigMart_Sales")
df2_dpp_partitioned.show(2)

df2_dpp_non_partitioned= df2.write.format("parquet").mode("append").option("path", "/content/dpp2_non_partitioned_BigMart_Sales").save()
df2_dpp_non_partitioned = spark.read.parquet("dpp2_non_partitioned_BigMart_Sales")
df2_dpp_non_partitioned.show(5)
```

Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Size	Outlet_Type
NULL	Regular	0.014753811	Fruits and Vegetables	231.7958	OUT027	1985	Medium	Outlet
20.35	Regular	0.0	Fruits and Vegetables	234.4958	OUT045	2002	NULL	Outlet

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 2 rows

Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year
FDA15	9.3	Low Fat	0.016047301	Dairy	249.8092	OUT049	1999
DRC01	5.92	Regular	0.019278216	Soft Drinks	48.2692	OUT018	2009
FDN15	17.5	Low Fat	0.016760075	Meat	141.618	OUT049	1999
FDX07	19.2	Regular	0.0	Fruits and Vegeta...	182.095	OUT010	1998
NCD19	8.93	Low Fat	0.0	Household	53.8614	OUT013	1987

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows

```
!rm -r dpp_non_partitioned_BigMart_Sales
```

```
import os
os.listdir("/content")
!find /content -name "*dpp2_*"
!rm -r dpp2_partitioned_BigMart_Sales
!ls -R /content

➡ /content/dpp2_non_partitioned_BigMart_Sales
rm: cannot remove 'dpp2_partitioned_BigMart_Sales': No such file or directory
/content:
BigMart_Sales.csv                dpp_partitioned_BigMart_Sales
dpp2_non_partitioned_BigMart_Sales  sample_data
dpp_non_partitioned_BigMart_Sales

/content/dpp2_non_partitioned_BigMart_Sales:
part-00000-5ea3b8e9-5090-4f50-82e4-162b6c3c7a1a-c000.snappy.parquet  _SUCCESS

/content/dpp_non_partitioned_BigMart_Sales:
part-00000-aa477b3a-cccc-4bc7-87ca-9f4d0e40c424-c000.snappy.parquet  _SUCCESS

/content/dpp_partitioned_BigMart_Sales:
'Outlet_Type=Grocery Store'      'Outlet_Type=Supermarket Type2'  _SUCCESS
'Outlet_Type=Supermarket Type1'  'Outlet_Type=Supermarket Type3'

'/content/dpp_partitioned_BigMart_Sales/Outlet_Type=Grocery Store':
part-00000-f41d5f25-07eb-4d91-9f01-8115194828b4.c000.snappy.parquet

'/content/dpp_partitioned_BigMart_Sales/Outlet_Type=Supermarket Type1':
part-00000-f41d5f25-07eb-4d91-9f01-8115194828b4.c000.snappy.parquet

'/content/dpp_partitioned_BigMart_Sales/Outlet_Type=Supermarket Type2':
part-00000-f41d5f25-07eb-4d91-9f01-8115194828b4.c000.snappy.parquet

'/content/dpp_partitioned_BigMart_Sales/Outlet_Type=Supermarket Type3':
part-00000-f41d5f25-07eb-4d91-9f01-8115194828b4.c000.snappy.parquet

/content/sample_data:
anscombe.json          mnist_test.csv
california_housing_test.csv  mnist_train_small.csv
california_housing_train.csv  README.md
```

Start coding or [generate](#) with AI.

```
df2_joined = df2_dpp_partitioned.join(df2_dpp_non_partitioned.filter(col("Outlet_Type")=="Grocery Store"), on="Item_Identifier", how="inner")
df2_joined.show()
```

➡

Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_S
DRA24	19.35	Regular	0.040154087	Soft Drinks	164.6868	OUT017	2007	N
DRA24	19.35	Regular	0.040154087	Soft Drinks	164.6868	OUT017	2007	N
DRA24	19.35	Regular	0.039920687	Soft Drinks	163.3868	OUT035	2004	Sn
DRA24	19.35	Regular	0.039920687	Soft Drinks	163.3868	OUT035	2004	Sn
DRA24	19.35	Regular	0.039990314	Soft Drinks	165.0868	OUT049	1999	Med
DRA24	19.35	Regular	0.039990314	Soft Drinks	165.0868	OUT049	1999	Med
DRA24	19.35	Regular	0.039895009	Soft Drinks	162.4868	OUT013	1987	F
DRA24	19.35	Regular	0.039895009	Soft Drinks	162.4868	OUT013	1987	F
DRA24	NULL	Regular	0.069909188	Soft Drinks	163.2868	OUT019	1985	Sn
DRA24	NULL	Regular	0.069909188	Soft Drinks	163.2868	OUT019	1985	Sn
DRA24	19.35	Regular	0.066831682	Soft Drinks	163.8868	OUT010	1998	N
DRA24	19.35	Regular	0.066831682	Soft Drinks	163.8868	OUT010	1998	N
DRA24	NULL	Regular	0.039734882	Soft Drinks	165.7868	OUT027	1985	Med

	DRA24	NULL	Regular	0.039734882	Soft Drinks	165.7868	OUT027	1985	Med
	FD011	8.0	Regular	0.030311951	Breads	247.4092	OUT049	1999	Med
	FD011	8.0	Regular	0.030264897	Breads	250.3092	OUT046	1997	Sr
	FD011	8.0	Regular	0.050657232	Breads	249.9092	OUT010	1998	N
	FD011	NULL	Regular	0.030118338	Breads	248.8092	OUT027	1985	Med
	FDU24	6.78	Regular	0.140955857	Baking Goods	92.212	OUT017	2007	N
	FDU24	6.78	Regular	0.0	Baking Goods	94.012	OUT013	1987	T
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
only showing top 20 rows									

```
df2_joined.explain(mode="formatted")
```



(1) Scan parquet
Output [12]: [Item_Identifier#1366, Item_Weight#1367, Item_Fat_Content#1368, Item_Visibility#1369, Item_Type#1370, Item_MRP#1371, Outl
Batched: true
Location: InMemoryFileIndex [file:/content/dpp2_partitioned_BigMart_Sales]
PushedFilters: [IsNotNull(Item_Identifier)]
ReadSchema: struct<Item_Identifier:string,Item_Weight:double,Item_Fat_Content:string,Item_Visibility:double,Item_Type:string,Item_MRP:

(2) ColumnarToRow [codegen id : 1]
Input [12]: [Item_Identifier#1366, Item_Weight#1367, Item_Fat_Content#1368, Item_Visibility#1369, Item_Type#1370, Item_MRP#1371, Outle

(3) Filter [codegen id : 1]
Input [12]: [Item_Identifier#1366, Item_Weight#1367, Item_Fat_Content#1368, Item_Visibility#1369, Item_Type#1370, Item_MRP#1371, Outle
Condition : isnotnull(Item_Identifier#1366)

(4) Exchange
Input [12]: [Item_Identifier#1366, Item_Weight#1367, Item_Fat_Content#1368, Item_Visibility#1369, Item_Type#1370, Item_MRP#1371, Outle
Arguments: hashpartitioning(Item_Identifier#1366, 200), ENSURE_REQUIREMENTS, [plan_id=881]

(5) Sort [codegen id : 2]
Input [12]: [Item_Identifier#1366, Item_Weight#1367, Item_Fat_Content#1368, Item_Visibility#1369, Item_Type#1370, Item_MRP#1371, Outle
Arguments: [Item_Identifier#1366 ASC NULLS FIRST], false, 0

(6) Scan parquet
Output [12]: [Item_Identifier#1464, Item_Weight#1465, Item_Fat_Content#1466, Item_Visibility#1467, Item_Type#1468, Item_MRP#1469, Outl
Batched: true
Location: InMemoryFileIndex [file:/content/dpp2_non_partitioned_BigMart_Sales]
PushedFilters: [IsNotNull(Outlet_Type), EqualTo(Outlet_Type,Grocery Store), IsNotNull(Item_Identifier)]
ReadSchema: struct<Item_Identifier:string,Item_Weight:double,Item_Fat_Content:string,Item_Visibility:double,Item_Type:string,Item_MRP:

(7) ColumnarToRow [codegen id : 3]
Input [12]: [Item_Identifier#1464, Item_Weight#1465, Item_Fat_Content#1466, Item_Visibility#1467, Item_Type#1468, Item_MRP#1469, Outle

(8) Filter [codegen id : 3]
Input [12]: [Item_Identifier#1464, Item_Weight#1465, Item_Fat_Content#1466, Item_Visibility#1467, Item_Type#1468, Item_MRP#1469, Outle
Condition : ((isnotnull(Outlet_Type#1474) AND (Outlet_Type#1474 = Grocery Store)) AND isnotnull(Item_Identifier#1464))

(9) Exchange
Input [12]: [Item_Identifier#1464, Item_Weight#1465, Item_Fat_Content#1466, Item_Visibility#1467, Item_Type#1468, Item_MRP#1469, Outle
Arguments: hashpartitioning(Item_Identifier#1464, 200), ENSURE_REQUIREMENTS, [plan_id=890]

(10) Sort [codegen id : 4]
Input [12]: [Item_Identifier#1464, Item_Weight#1465, Item_Fat_Content#1466, Item_Visibility#1467, Item_Type#1468, Item_MRP#1469, Outle
Arguments: [Item_Identifier#1464 ASC NULLS FIRST], false, 0

(11) SortMergeJoin [codegen id : 5]
Left keys [1]: [Item_Identifier#1366]
Right keys [1]: [Item_Identifier#1464]
Join type: Inner
Join condition: None

(12) Project [codegen id : 5]
Output [23]: [Item_Identifier#1366, Item_Weight#1367, Item_Fat_Content#1368, Item_Visibility#1369, Item_Type#1370, Item_MRP#1371, Outl
Input [24]: [Item_Identifier#1366, Item_Weight#1367, Item_Fat_Content#1368, Item_Visibility#1369, Item_Type#1370, Item_MRP#1371, Outle

Start coding or [generate](#) with AI.

With AQE

```
spark.conf.set("spark.sql.adaptive.enabled", "true")
spark.conf.get("spark.sql.adaptive.enabled")
```

```
'true'
```

```
df = spark.read.format("csv") \
    .option("inferSchema", True) \
    .option("header", True) \
    .load("BigMart_Sales.csv")
```

```
df.show(2)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|Item_Identifier|Item_Weight|Item_Fat_Content|Item_Visibility| Item_Type|Item_MRP|Outlet_Identifier|Outlet_Establishment_Year|Outlet_Si
|-----+-----+-----+-----+-----+-----+-----+-----+-----+
|      FDA15|      9.3|      Low Fat|      0.016047301|      Dairy|249.8092|      OUT049|      1999|      Medi
|      DRC01|      5.92|      Regular|      0.019278216|Soft Drinks| 48.2692|      OUT018|      2009|      Medi
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 2 rows
```

```
df.rdd.getNumPartitions()
```

```
1
```

```
df_new_with_AQE = df.groupBy("Item_Fat_Content").count()
df_new_with_AQE.show()
```

```
+-----+-----+
|Item_Fat_Content|count|
+-----+-----+
|      low fat|    112|
|      Low Fat|   5089|
|      LF|     316|
|      Regular|  2889|
|      reg|     117|
+-----+-----+
```

```
df_new_with_AQE.explain(mode="formatted")
```

```
== Physical Plan ==
AdaptiveSparkPlan (5)
+- HashAggregate (4)
   +- Exchange (3)
      +- HashAggregate (2)
         +- Scan csv (1)

(1) Scan csv
Output [1]: [Item_Fat_Content#289]
Batched: false
Location: InMemoryFileIndex [file:/content/BigMart_Sales.csv]
ReadSchema: struct<Item_Fat_Content:string>

(2) HashAggregate
Input [1]: [Item_Fat_Content#289]
Keys [1]: [Item_Fat_Content#289]
Functions [1]: [partial_count(1)]
Aggregate Attributes [1]: [count#435L]
Results [2]: [Item_Fat_Content#289, count#436L]

(3) Exchange
Input [2]: [Item_Fat_Content#289, count#436L]
Arguments: hashpartitioning(Item_Fat_Content#289, 200), ENSURE_REQUIREMENTS, [plan_id=262]

(4) HashAggregate
Input [2]: [Item_Fat_Content#289, count#436L]
Keys [1]: [Item_Fat_Content#289]
Functions [1]: [count(1)]
Aggregate Attributes [1]: [count(1)#425L]
Results [2]: [Item_Fat_Content#289, count(1)#425L AS count#426L]

(5) AdaptiveSparkPlan
Output [2]: [Item_Fat_Content#289, count#426L]
Arguments: isFinalPlan=false
```


Start coding or [generate](#) with AI.