

## Dynamic Partition Pruning (DPP)

### What is Partition Pruning?

In partitioned tables, data is split into **subdirectories** (partitions) based on partition keys (e.g., date, region).

**Partition Pruning** means reading **only the partitions needed** for a query instead of scanning the entire dataset.

This reduces I/O and speeds up queries.

### Static vs Dynamic Partition Pruning

Feature	Static Partition Pruning	Dynamic Partition Pruning
Partition filter value	Known at <b>query compile time</b>	Determined <b>at runtime</b>
Example	WHERE region = 'US'	WHERE region IN (SELECT region FROM other_table)
When used	Simple filters	Filters depend on another query result or runtime parameter
Performance	Very efficient	Efficient but slightly more overhead due to runtime evaluation

### Why Dynamic Partition Pruning is Needed

Many queries have partition filter values **not known in advance**.

```
SELECT * FROM sales WHERE region IN ( SELECT region FROM active_regions WHERE year = 2024 );
```

The region values come from another table.

Without **DPP**, all partitions of sales would be scanned → slow and expensive.

With **DPP**, only partitions matching the runtime region values are read.

### How DPP Works in Spark

**Static PP** happens at compile/optimization time.

**Dynamic PP:**

1. Spark compiles the query with a placeholder for the partition filter.
2. At runtime, executes the subquery to get filter values.
3. Applies partition pruning before scanning the table.

```
SET spark.sql.optimizer.dynamicPartitionPruning.enabled=true;
```

```
SELECT * FROM fact_sales WHERE date_id IN (  
  SELECT DISTINCT date_id FROM dim_calendar WHERE holiday_flag = 'Y' );
```

Spark Configs to Enable

```
SET spark.sql.optimizer.dynamicPartitionPruning.enabled = true;
```

This tells Spark to actually perform runtime partition pruning.

## Visual Example

Imagine your dataset has partitions:

/sales/region=US/

/sales/region=CA/

/sales/region=UK/

/sales/region=IN/

Query asks for:

```
SELECT * FROM sales WHERE region IN (SELECT region FROM active_regions);
```

Without DPP: Scans all 4 partitions.

With DPP: Scans only the partitions returned by active\_regions (maybe just US and CA).

## How DPP Works in Spark

### 1. Compilation Phase

- Spark creates a plan with a placeholder for the partition filter.
- Knows it will get the filter values from another query stage.

### 2. Runtime

- Executes the filter-producing subquery first.
- Obtains actual values (e.g., ['US', 'CA']).
- Applies these values to the scan of the partitioned table.

### 3. Scan

- Reads only the relevant partitions (e.g., /region=US/ and /region=CA/).
- Skips all others.

## Visual Example

Imagine your dataset has partitions:

/sales/region=US/

/sales/region=CA/

/sales/region=UK/

/sales/region=IN/

```
SET spark.sql.optimizer.dynamicPartitionPruning.enabled = true;
```

```
SELECT * FROM sales WHERE region IN ( SELECT region FROM active_regions WHERE year = 2025);
```

## Without DPP

- Spark reads all four partitions.
- Filters out unwanted rows after reading.

## With DPP

- Spark first runs `SELECT region FROM active_regions WHERE year=2025`
- Suppose result = ['US', 'CA']
- Spark reads only /region=US/ and /region=CA/.

## Before vs After Diagram

Before DPP (Full Scan)

Query Plan:

```
| Scan all partitions in sales |
```

|

Filter region from subquery

|

Final Output

Reads **US, CA, UK, IN** partitions → **Slow**.

After DPP (Pruned Scan)

Step 1: Run subquery

```
| SELECT region FROM active_regions |  
| Result: US, CA                    |
```

|

Step 2: Prune partitions

```
| Scan only region=US, CA          |
```

|

Final Output

Reads **only needed partitions** → **Fast**.

## Benefits

- **Performance:** Avoids unnecessary I/O.
- **Cost Savings:** In cloud environments (S3, GCS, ADLS), reading less data = lower cost.
- **Scalability:** Handles large datasets efficiently.

## Spark Configs for DPP

-- Enable Dynamic Partition Pruning

```
SET spark.sql.optimizer.dynamicPartitionPruning.enabled = true;
```

-- Optional: Increase timeout if filter value fetch takes longer

```
SET spark.sql.dynamicPartitionPruning.broadcast.timeout = 5m;
```

## Key Takeaways

- Use DPP when your filter values come from another dataset or runtime parameter.
- Ensure your table is partitioned **on the filter column** for maximum effect.
- Validate in **Spark UI** by checking the number of partitions scanned.